

MYSORE SRIDHAR

AI Engineer — Generative AI & Agentic AI

☎ +91-7396947527

✉ mysorensridhar72@gmail.com

🌐 sridhar-mysore.vercel.app

🔗 github.com/sapphire2207

🌐 linkedin.com/in/mysorensridhar

📍 Hyderabad

EDUCATION

B.Tech Computer Science in Artificial Intelligence and Data Science

2022–2026

Sri Indu College of Engineering & Technology

CGPA: 8.59/10

Intermediate Education (MPC)

Chaithanya Kalasala Govt Aided Jr College

GPA: 9.3/10

SSC

Bhashyam High School

GPA: 10.0/10

SKILLS

Python, LangChain, LangGraph, Generative AI, Agentic AI, LLMs, Prompt Engineering, RAG, AI Agents, Multi-Agent Systems, FastAPI, Pydantic, Vector Databases, ChromaDB, Embeddings, AI Integrations, REST APIs, Git, Next.js, TypeScript

PROJECTS

ResearchFlow AI – Multi-Agent AI Research Platform

<https://github.com/sapphire2207/research-flow-ai>

A production-ready multi-agent AI research platform that autonomously searches, analyzes, synthesizes, and critiques web content to generate comprehensive research reports.

- Orchestrated a production-grade multi-agent research platform using LangGraph, LangChain, FastAPI, Pydantic, Next.js, and TypeScript, coordinating Search, Reader, Writer, and Critic AI agents through a stateful graph-based workflow.
- Implemented real-time research execution with Server-Sent Events (SSE), live progress streaming, and intelligent fallback mechanisms, reducing perceived response latency by 50% while improving workflow reliability by 40%.
- Integrated autonomous web intelligence by combining Tavily Search, BeautifulSoup, and Mistral AI to retrieve, analyze, validate, and synthesize information into structured research reports with automated quality scoring.
- Architected a modular AI backend featuring reusable LangChain components, strongly typed Pydantic schemas, and scalable REST APIs, enabling autonomous research across 20+ production-ready frontend and backend modules.

InkFlow AI – Multi-Agent AI Blog Generation Platform

<https://github.com/sapphire2207/inkflow-ai>

A planning-based multi-agent AI content generation platform that researches, writes, and enriches publication-ready blogs with citations and AI-generated diagrams.

- Engineered a planning-driven multi-agent content generation platform using LangGraph, LangChain, FastAPI, Pydantic, Next.js, and TypeScript, orchestrating Router, Research, Orchestrator, Worker, and Reducer AI agents.
- Automated evidence-backed content creation by integrating Tavily Search and Mistral AI, enabling parallel research, source attribution, and recency-aware information retrieval while reducing manual research effort by 75%.
- Optimized live AI workflow execution through Server-Sent Events (SSE), three-tier streaming recovery, and REST fallback strategies, increasing generation reliability by 40% across six processing stages.
- Enabled AI-powered technical documentation by integrating Pollinations AI to automatically generate and embed up to 3 contextual diagrams into publication-ready Markdown articles, reducing manual documentation effort by 60%.

Recap AI – AI Meeting Intelligence Platform

<https://github.com/sapphire2207/recap-ai>

An AI-powered meeting intelligence platform that transcribes, summarizes, extracts actionable insights, and enables conversational Q&A over meeting recordings using RAG.

- Engineered an enterprise-grade meeting intelligence platform using LangGraph, LangChain, FastAPI, Pydantic, Next.js, and TypeScript, automating multilingual transcription, executive insights, and conversational meeting Q&A through a parallel DAG architecture.
- Streamlined multilingual speech processing by integrating OpenAI Whisper, Sarvam AI, and automated audio chunking, improving transcription robustness for English and Hinglish recordings while reducing processing failures by 50%.
- Accelerated AI-driven meeting analysis by executing concurrent LangGraph pipelines that generated meeting titles, executive summaries, action items, key decisions, and open questions, reducing end-to-end latency by 45%.
- Implemented a transcript-aware Retrieval-Augmented Generation (RAG) pipeline using ChromaDB, HuggingFace Embeddings, LangChain LCEL, and Mistral AI, enabling persistent vector indexing and context-grounded conversational search.

KEY ACHIEVEMENTS

SMART INDIA HACKATHON

- Worked with a team of 4 in the Railway Complaint Management project, where we delivered a presentation that secured us a spot in the 2nd round out of 25 teams.

College Website

- Collaborated with a team of 2 to create a dynamic college website for the Artificial Intelligence and Data Science club, featuring announcements, syllabus, and curriculum.